
Separating the Art from the Algorithm: Detecting Moral Decoupling in Consumer Discourse About Generative AI

Onuchukwu Joseph Chimezie¹, Julius Mboli¹

¹Centre of Excellence for Data Science, Artificial Intelligence
and Modelling (DAIM), University of Hull, United Kingdom
prolifiquemexie@gmail.com, mboli4god@gmail.com

Abstract: Some consumers who learn that a generative AI tool was trained on creative work taken without consent condemn the practice and stop using it. Others condemn it and keep using it. The second response is moral decoupling: separating judgements about wrongdoing from judgements about usefulness. The construct is established in laboratory research but has not been measured in naturally occurring discourse. We present a stance detection study of 7,371 Reddit posts from three professional communities differently exposed to generative AI. Each post is human-annotated on two axes, stance toward AI content and type of justification, with decoupling operationalised as a supportive stance co-occurring with explicit ethical acknowledgement. Across five calibration rounds two annotators reached Cohen’s kappa of 0.654 on stance and 0.621 on argument type; GPT-4o, given the same guidelines, reached 0.058 and 0.209 and never used the MIXED label. Fine-tuned RoBERTa reached macro-F1 of 0.63 and 0.82. Moral positions differ sharply across communities ($V = 0.23$), and among engaged posts the decoupling-coupling balance differs too ($V = 0.15$). Marketers rarely engage morally at all, whereas artists produce most moral reasoning and resolve it against the tools three times in four. Involvement governs primarily whether ethical evaluation happens, secondarily which way it falls.

Keywords: moral decoupling, stance detection, generative AI, AI ethics, computational social science, natural language processing

1. Introduction

Generative AI tools have been adopted faster than almost any prior consumer technology, and they have drawn ethical criticism for almost as long. Image generators such as Stable Diffusion and Midjourney, and text generators such as the GPT family, were trained on enormous collections of human work gathered largely without consent. Artists have been the most vocal: a recent survey of 459 practising artists found that a clear majority believe their work should not be used in training without disclosure, and worry about both their livelihoods and who profits from models built on their output [1]. Those concerns are not idle, since the same generation of tools is now estimated to expose a large share of creative and knowledge work to automation [2]. And yet adoption keeps climbing, often among the very professionals voicing the loudest objections. That gap, between what people say about how these tools were built and what they actually do with them, is the puzzle this paper takes up.

Consumer psychology offers a precise name for one way that gap is bridged. Bhattacharjee et al. [3] describe moral decoupling as a strategy in which a person separates the judgement of how wrong an act is from the judgement of how useful the actor or product is. It differs from rationalisation in a way that matters for measurement: rationalisation lowers the perceived wrongdoing, while decoupling leaves it intact and sets

it aside. A user can hold that training on scraped work is wrong and still use the tool, never reconciling the two. The strategy is well documented for transgressing celebrities and brands, but almost always through constructed scenarios [4, 5]. What remains open is whether the same reasoning surfaces in the unprompted talk of real users, and whether its prevalence tracks how close a person stands to the harm.

This paper makes three contributions. First, we give moral decoupling a measurable textual definition, as a supportive stance toward AI tools combined with an explicit deontological acknowledgement that their construction was ethically problematic. Second, we build an annotation scheme and a human-labelled corpus of 7,371 posts from three communities with differing exposure to generative AI. Third, we test whether decoupling tracks involvement, and report a result that overturns the obvious prediction and replaces it with a sharper one: the difference between communities lies first in whether they treat the matter as a moral question at all, and only secondarily in which side they take once they do.

2. Related Work

2.1. Moral decoupling

Moral decoupling was introduced to explain why consumers keep supporting athletes, celebrities and companies who have transgressed [3]: a preserved moral judgement is separated from the performance judgement that drives behaviour. A recurring theme in later work on corporate misconduct is that involvement with the actor shapes how strongly the strategy is used [4], which is the premise behind the comparative design here. The construct has now reached generative AI directly. Liu et al. [5] ran two experiments on 509 respondents and found moral decoupling carrying the effect of GenAI misuse on brand cancellation. That study is the closest precedent for ours, and the contrast is instructive: it manipulates a transgression and measures the reasoning that follows, whereas we ask whether such reasoning arises at all when nobody is prompting it. The artist survey of Lovato et al. [1] reinforces the lens, showing that those most exposed reason in exactly the utility-versus-rights terms our scheme captures.

The construct also needs distinguishing from *moral disengagement*, in which the judgement is neutralised in advance through mechanisms such as diffusion of responsibility or sanitising language [6]. Decoupling is the narrower case: the judgement of wrongdoing survives intact and is held apart from the judgement of usefulness. Much of what we label pure utilitarian discourse may be closer to disengagement, or to the question never arising, and our scheme does not separate those possibilities.

2.2. Stance detection and split targets

Identifying whether a text supports or opposes a target is stance detection, which differs from sentiment analysis because polarity is defined relative to a specific object rather than to overall tone [7]. Decoupling breaks ordinary sentiment tools because its target is split. Consider two sentences with near-identical surface negativity: “this is theft, but I still use it” and “this is theft, so I refuse.” They express opposite behavioural positions. Recovering decoupling therefore requires stance and the form of justification to be recorded as separate dimensions, an aspect-oriented framing rather than a document-level one.

2.3. Transformer models for social text

Fine-tuned transformer encoders are standard for classifying social text. RoBERTa [8] refined the BERT pre-training recipe and remains a strong baseline. Our posts often run to several paragraphs, which favours a model pre-trained on longer text over one tuned for very short messages; we use RoBERTa-base for that reason.

3. Data

3.1. Community selection

Three Reddit communities anchor the study, chosen because they sit at different points on a gradient of exposure to generative AI. r/marketing treats the tools mainly as instruments for producing content and stands at some remove from the ethical dispute, though the ethics of GenAI in content marketing have been mapped through a deontological lens [9]. r/ArtistLounge is made up of practising visual artists whose work may sit in the training data and who compete directly with image generators for commissions: the group most exposed to the harm [1]. r/copywriting falls between the two. Comparing along this gradient is what lets the involvement hypothesis be tested rather than asserted.

3.2. Collection and filtering

Posts and comments were retrieved with the Arctic Shift download tool¹, which returns posts and comments for a given subreddit over a specified date range. Six files were downloaded, posts and comments for each of the three communities, spanning January 2022 to May 2025, a window that brackets the public releases of DALL-E 2, Midjourney, Stable Diffusion and ChatGPT, and totalling 1,636,493 records. We kept items mentioning AI-related terms, then discarded deleted and removed content, bot output, posts scoring below two and posts under ten words, and de-duplicated on text. A stratified sample balanced the three communities against the volume available in each. The final corpus holds 7,371 posts; Table 1 summarises it. Because re-sampling a live archive is not bit-for-bit reproducible, the assembled corpus is frozen and shipped with the project as the single source of truth, and every downstream analysis reads only that file.

Table 1. Corpus composition by community.

Community	Posts	Mean words	Median words
r/ArtistLounge	3,000	128.3	89
r/marketing	2,500	142.7	71
r/copywriting	1,871	137.3	78
Total	7,371	n/a	n/a

3.3. Ethics and anonymisation

Only public posts were used. Usernames and URLs were stripped and each post given an internal numeric identifier, and the work was conducted under the University of Hull Faculty of Science and Engineering ethics process for research on public online data. Verbatim quotation carries a residual risk that a post, and with it its author, can be located through search even where usernames are removed; the illustrative quotations in Appendix A were chosen to be representative rather than distinctive, and none is attributed beyond its subreddit. Public availability does not amount to consent to research use, which is why identifiers are stripped and all reported findings are aggregate.

4. Annotation Scheme

Each post carries two primary labels and three supplementary ones. AI-STANCE records the position taken toward AI-generated content as PRO, ANTI, NEUTRAL or MIXED. ARG-TYPE records the dominant justification as UTILITARIAN, appealing to efficiency, cost or quality; DEONTOLOGICAL, appealing to rights, consent or fairness; MIXED; or UNCLEAR. The supplementary labels capture whether ethical acknowledgement is explicit or implicit, where blame is directed, and whether language about authenticity and the human element is present.

The construct of interest follows from the two primary labels. *Moral decoupling* is a PRO stance with DEONTOLOGICAL reasoning: support for the tool alongside acknowledgement of the ethical problem.

¹<https://arctic-shift.photon-reddit.com/download-tool>

Moral coupling is an ANTI stance with DEONTOLOGICAL reasoning, where the objection aligns with rejection. A PRO stance with UTILITARIAN reasoning and no ethical content is *pure utilitarian*, kept distinct from decoupling. That distinction does real work in the results: it is the difference between setting an objection aside and never raising one.

Two caveats attach to this operationalisation. A post that supports limited use while acknowledging an ethical problem may reflect ambivalence, reluctant adoption or pragmatic constraint rather than decoupling in the psychological sense, and a single post cannot establish that its author preserved a moral judgement and consciously set it aside. We therefore treat the PRO-plus-deontological combination as a *textual proxy* for decoupling, not as direct evidence of a cognitive strategy, and the claims that follow are claims about discourse rather than about minds. The gating is also strict: posts pairing a MIXED stance with deontological reasoning, plausibly authors in unresolved dissonance, fall outside the count, so the incidence reported below is a lower bound and a sensitivity analysis across that threshold would test the boundary.

4.1. Calibration and agreement

Two annotators independently labelled a pilot of 400 posts. Agreement began low, which is what one should expect when the task is to infer moral reasoning from short, informal text. We treated that as a starting point rather than a verdict. Across five calibration rounds we examined the confusion matrices after each pass, isolated the recurring sources of disagreement, rewrote the guidelines to address them, and re-annotated. Agreement on stance climbed from 0.189 to 0.654, and on argument type from 0.235 to 0.621 (Table 4).

The conventional threshold for substantial agreement is 0.61, so both primary dimensions clear it, stance comfortably and argument type narrowly. The three supplementary dimensions improved over the first four rounds and then stopped, flat across rounds four and five. That plateau is informative in itself: the limit on those categories is how they are defined, not how well the annotators are trained, so we report them descriptively and rest the main analysis on the two primary dimensions. Remaining disagreements were settled by adjudication against the written guidelines rather than by negotiated consensus, so no label reflects one coder’s view prevailing over another’s.

Once agreement was established, the two annotators divided the remaining posts between them and labelled them using the calibrated guidelines, so all 7,371 posts carry human labels on both primary dimensions. Reliability is measured on the 400-post double-coded pilot; posts beyond it were single-coded, with the same adjudication rule applied to any case a coder flagged as uncertain. The classifiers reported in Section 5 are evaluated against these human labels and are not used to generate any of the distributions, tests or prevalence figures reported in Section 6; no predicted label enters the substantive analysis.

4.2. Zero-shot LLM annotation comparison

We also asked whether a large language model could take over the labelling. GPT-4o was given the same written guidelines and applied to the 400-post pilot under two prompt designs, both zero-shot, with no few-shot exemplars and no chain-of-thought scaffolding. Both failed, in opposite directions. The first collapsed toward the safe middle, assigning NEUTRAL to 368 of 400 posts, and reached Cohen’s kappa of 0.025 on stance and 0.110 on argument type. A re-engineered prompt over-committed to PRO (210 of 400), reaching 0.058 and 0.209, while explicitness and blame-target agreement fell from 0.130 to 0.066 and 0.076 to 0.037: prompt engineering traded one failure mode for another.

Two patterns are worth isolating. Neither run assigned the MIXED stance label even once, where the human coders used it for 94 of the 400 posts, so the model could not represent the ambivalence on which the decoupling construct depends. And agreement on the authenticity dimension was identical in both runs at 0.512, far above any interpretive dimension, which is unsurprising because that category approximates keyword matching. We report this as a methodological result: automated annotation of moral reasoning was not reliable here, and human judgement was required. Because only zero-shot prompting was evaluated, this establishes that this regime fails rather than that the task lies beyond current models; few-shot and chain-of-thought variants remain untested.

5. Classification Models

The corpus was split 70/15/15 into training, validation and test sets, stratified by stance and community. We evaluate two model families on the two primary tasks and report macro-averaged F1, the honest metric under the class imbalance these labels carry.

5.1. Baselines

The baseline is a linear support vector machine over TF-IDF unigrams and bigrams. To test whether an explicit moral lexicon adds anything, we also evaluate a variant that concatenates the sparse TF-IDF matrix with a small dense set of normalised moral-category frequencies drawn from a standard psycholinguistic dictionary [10]. It made no difference: macro-F1 moved by less than 0.002 on both tasks, and downward on argument type. The lexical features had already absorbed the moral vocabulary, and the dictionary contributed no measurable signal. That null is worth stating rather than burying, since it cautions against assuming hand-built moral lexicons add value once n-grams are present.

5.2. Fine-tuned RoBERTa

We fine-tune RoBERTa-base [8] separately for each task with a standard classification head and early stopping on validation macro-F1. RoBERTa beats the baseline on both tasks, and on argument type the margin is large (Table 2). Training used a learning rate of 2×10^{-5} , batch size 16, a maximum sequence length of 256 tokens, linear warmup, and early stopping with a patience of two epochs over a maximum of five. Splits were stratified by stance and community with a fixed random seed.

Table 2. Test-set macro-F1 by model and task.

Model	AI-STANCE	ARG-TYPE
TF-IDF + SVM	0.532	0.600
TF-IDF + moral features	0.532	0.598
RoBERTa-base	0.629	0.823

On stance, RoBERTa reaches 0.69 accuracy, handling NEUTRAL and ANTI most reliably and the rare MIXED class least. On argument type it reaches 0.88 accuracy, strong on UTILITARIAN, DEONTOLOGICAL and UNCLEAR and weakest on MIXED. The errors are concentrated where the categories genuinely overlap: on argument type the MIXED class is the weakest by a wide margin and bleeds primarily into UTILITARIAN, while the linear baseline recovers only 2 of 89 cases, a recall of 0.02. The same categories that troubled the human annotators trouble the model, which supports reading the difficulty as a property of the construct rather than of any one classifier. A confusion analysis (Appendix C) shows the errors run along the reasoning axis rather than the stance axis: 84 of the 240 posts carrying moral reasoning are misread as purely utilitarian, while only 6 of 866 non-moral posts are labelled moral. Decoupling is therefore confused with pure utilitarianism, not with coupling, and the bias is conservative.

6. Results

6.1. Distribution of moral positions

Table 3 gives the distribution of moral positions by community, and two things are hard to miss. The decoupling pattern is rare everywhere, and it is in fact slightly more common in r/ArtistLounge (3.5%) than in r/marketing (0.8%) or r/copywriting (1.0%), which is the reverse of what the involvement prediction would suggest. The larger story is the gulf in how much moral reasoning of any kind each community produces. In ArtistLounge, coupling alone reaches 10.6% and mixed moral reasoning 13.2%; in marketing those same categories sit at 1.0% and 4.0%, and the discourse is instead dominated by pure utilitarian content at 24.8%.

Table 3. Moral position distribution by community (% of community posts).

Position	ArtistLounge	marketing	copywriting
Moral decoupling	3.5	0.8	1.0
Moral coupling	10.6	1.0	1.3
Mixed reasoning	13.2	4.0	3.9
Pure utilitarian	10.8	24.8	21.4
Practical rejection	8.1	7.0	11.3
Other / none	53.8	62.4	61.1

6.2. Hypothesis tests

A chi-square test of independence between community and moral position is decisive: $\chi^2 = 782.57$ ($df = 10$, $p < 0.001$), with Cramér’s V of 0.230, a substantial association for discourse data. The communities clearly differ. We then narrow to the 511 posts that are either decoupling or coupling and ask whether the balance between them depends on community. It does: $\chi^2 = 11.12$ ($df = 2$, $p = 0.004$), Cramér’s V = 0.147. Within r/ArtistLounge, coupling outnumbers decoupling by roughly three to one (319 against 104, so 24.6% of rights-based reasoning ends in endorsement), whereas the engaged posts in the two commercial communities divide close to evenly (41.9% in r/copywriting and 42.2% in r/marketing).

The commercial cells are small (43 and 45 posts), so the test was repeated by Monte Carlo permutation (200,000 replicates, $p = 0.004$) and by Fisher exact tests, which reproduce it: r/ArtistLounge differs from r/copywriting ($p = 0.018$) and from r/marketing ($p = 0.019$), while the two commercial communities are statistically indistinguishable from each other ($p = 1.00$). Pooled against r/ArtistLounge they give an odds ratio of 0.45 (Fisher $p = 0.002$, minimum expected count 24.3). The direction therefore does not rest on sparse cells, though the effect is modest and we report it as a refinement rather than a settled result.

That pair of results is the heart of the paper, and the two levels must be kept apart because they condition on different denominators. Across all posts, the dominant effect by some margin is whether moral reasoning occurs at all: artists produce it at roughly four times the rate of either commercial community, and it is this gap that carries the large overall association. Within the minority who do reason about rights, direction also varies, and in the opposite direction to the raw rates: a far greater share of engaged commercial users resolve the question in favour of continued use. Involvement therefore governs entry into moral evaluation first and foremost, and shapes its direction secondarily. The decoupling pattern is rare in marketing discourse because, for the most part, marketers never frame tool use as a moral question at all.

6.3. Distinctive vocabulary

A multinomial logistic regression over TF-IDF features, fitted to predict community, recovers the point lexically. ArtistLounge is marked by *art*, *artists*, *artist* and, tellingly, *copyright*, the only community with an ethics term among its most distinctive features. Copywriting is marked by *copywriting*, *copy* and client-facing words; marketing by *marketing*, *content*, *SEO* and tool vocabulary, with no ethics term in its leading set. The moral vocabulary is not outweighed in marketing discourse so much as absent from it, which is consistent with the survey finding that sustained ethical engagement clusters among artists themselves [1].

7. Discussion

We set out to test whether decoupling is more common among less-involved users, and found something more useful than a yes or no. The pattern is rare across the board, and the difference that matters most between communities is whether moral reasoning occurs at all. This reframes the involvement account rather than refuting it.

Three points of contrast with prior work clarify what the present study adds. First, against the experimental tradition that established moral decoupling, our contribution is one of setting rather than mechanism. Bhattacharjee et al. [3] and the re-examination by Haberstroh et al. [4] demonstrated decoupling under controlled conditions, presenting participants with a transgression and measuring how they reconcile admiration

with disapproval. That design guarantees moral engagement by construction: the scenario puts the wrongdoing in front of the participant. Our corpus removes that guarantee. When people are left to raise or ignore the ethical dimension themselves, most of the low-involvement group never raise it, and the decoupling pattern, which presupposes engagement, becomes correspondingly rare. We therefore observe a prior step that scenario-based work cannot, namely the decision whether to moralise at all, and we locate the involvement effect there first.

Second, against survey work in this domain, our findings agree on substance while differing in reach. Lovato et al. [1] surveyed 459 artists and found predominantly rights-based reasoning about consent, ownership and fair compensation; our ArtistLounge results echo that from unsolicited discourse rather than elicited responses. But a survey of artists cannot speak to the communities that are *not* surveyed, whereas our comparative design shows that the sustained ethical argument it documents is largely confined to the most exposed community.

Third, against the methodological grain of much recent text classification, we report two deliberately negative results that are easy to omit and useful to keep. A psycholinguistic moral lexicon [10] added no measurable signal over lexical features: the n-grams had already absorbed the moral vocabulary. And a zero-shot language model annotated the moral dimensions at close to chance under two prompt designs, while scoring 0.512 on the one dimension that reduces to keyword matching. The two results point the same way. Hand-built lexicons and general-purpose models both recover the lexical surface of moral discourse and both miss the reasoning underneath it, which suggests the limiting factor is the construct rather than the apparatus.

The practical implication follows directly. Ethical objections to generative AI are usually addressed to users as if the task were to win an argument those users are already having. If, for the least exposed groups, the argument is not taking place, moral persuasion has little to grip. The leverage point sits earlier, in whether the ethical dimension becomes salient at all, which is a question of agenda-setting rather than of argument.

8. Conclusions

To our knowledge this is the first computational measurement of moral decoupling in real discourse about generative AI, with an annotation scheme, a human-labelled corpus of 7,371 posts, and classifiers reaching macro-F1 of 0.63 on stance and 0.82 on argument type. The central result is that domain involvement governs whether consumers morally evaluate generative AI at all, and that this effect is considerably larger than its influence on the direction that evaluation takes, though direction varies too among the engaged minority. The group most exposed to the harm produces the most moral reasoning, for and against; the least exposed largely bypasses the question in favour of practical concerns.

8.1. Limitations

Several limitations bound these conclusions; Appendix E sets them out in full. Four matter most. The construct is a textual proxy: a supportive post acknowledging an ethical problem may reflect ambivalence or reluctant adoption, and the gating is strict enough that MIXED-stance posts with deontological reasoning fall outside the count, making 1.9 per cent a lower bound. The communities’ norms shape not only how moral reasoning is expressed but whether it appears at all, so the engagement gap may partly reflect platform governance and social desirability rather than private cognition. The involvement gradient is assumed from occupational composition rather than measured, and the communities differ in demographics, moderation and incentives as well. And posts are treated as independent although multiple posts may share an author, which we cannot quantify after anonymisation.

8.2. Recommendations for future work

The most informative extension is a cross-domain transfer test. Because both model families were trained on the pooled corpus, and because the distinctive-vocabulary analysis places *copyright* among the strongest

r/ArtistLounge discriminators, the classifiers may be exploiting community register rather than moral-reasoning structure; training on one community and testing on the others would separate the two. Appendix F lists four further directions: comparative breadth across platforms, extending the language model comparison beyond the pilot and beyond zero-shot prompting, a sensitivity analysis across the decoupling threshold, and an intervention study on whether making the ethical question salient changes behaviour.

Reproducibility

The analysis is organised as four ordered notebooks that run from the raw archives to the reported hypothesis tests. The corpus is frozen to a single CSV that every downstream notebook reads, so all figures reproduce exactly and independently of any later change to the live archive. All splits use a fixed random seed and are stratified by stance and community. The annotation guidelines issued to both annotators, and the label files used for the zero-shot comparison, are released with the code. All 7,371 posts carry human labels; no predicted label enters any reported distribution or test. Code and data: <https://doi.org/10.5281/zenodo.21891837>

Acknowledgements

The authors thank the second annotator for their contribution to the calibration rounds and to annotation of the corpus.

A. Annotation examples

Three short examples from the corpus, lightly de-identified, fix the categories. A copywriter writes:

“I used to hate AI and think it will steal my job. Now I love using them for research and brainstorming.”

The prior objection, that the tools threaten livelihoods, is named and then set aside in favour of usefulness; this is the decoupling pattern. By contrast, an artist couples the objection to rejection, steering others away on ethical grounds:

“You don’t have to use ChatGPT for it. There are other AI tools that aren’t on the unethical side of it.”

And a marketer shows the pure utilitarian pattern that dominates that community, where the ethical question never enters:

“ChatGPT is surprisingly helpful when I’m staring at a blank page. I use it all the time to get me started on copy.”

B. Calibration detail

Table 4 gives the full agreement trajectory across the five calibration rounds.

Table 4. Inter-annotator agreement (Cohen’s kappa) across five calibration rounds.

Dimension	R1	R2	R3	R4	R5
AI-STANCE	0.189	0.511	0.587	0.623	0.654
ARG-TYPE	0.235	0.391	0.505	0.589	0.621
Explicitness	0.206	0.369	0.383	0.383	0.383
Blame target	0.134	0.284	0.295	0.295	0.291
Authenticity	0.142	0.328	0.356	0.356	0.356

C. Confusion analysis

Because the decoupling construct is built from the argument-type axis, it matters where that axis fails. Table 5 gives the baseline confusion matrix for ARG-TYPE on the test set.

Table 5. ARG-TYPE confusion matrix, TF-IDF + SVM baseline. Rows are true labels, columns predicted.

True \ Predicted	DEON	MIXED	UNCL	UTIL
DEONTOLOGICAL	111	1	10	29
MIXED	31	2	1	55
UNCLEAR	4	0	218	74
UTILITARIAN	5	1	35	529

The errors are not symmetric, and their direction matters for the construct. UTILITARIAN is predicted 687 times against 570 true instances, while MIXED is predicted 4 times against 89. Of the 240 posts that carry moral reasoning of either kind, 84, or 35 per cent, are misread as purely utilitarian: 29 of 151 DEONTOLOGICAL posts and 55 of 89 MIXED ones. The remaining MIXED errors go the other way, 31 of 89 being read as fully deontological. In other words, the baseline resolves ambivalence by collapsing it, usually toward the practical pole, and only rarely invents moral reasoning where none exists: just 6 of 866 non-moral posts are labelled DEONTOLOGICAL or MIXED.

Two consequences follow. First, decoupling is not confused with coupling, since the errors run along the reasoning axis rather than the stance axis; it is confused with pure utilitarianism, which is precisely the distinction the scheme was built to preserve. Second, because the reported prevalence figures use human labels throughout, these errors do not propagate into Section 6. They do, however, bound what an automated replication could claim, and the direction of the bias is conservative: a pipeline of this kind would understate moral engagement rather than inflate it, and would understate decoupling in particular. Fine-tuned RoBERTa reduces the collapse substantially, lifting MIXED F1 from 0.04 to 0.66, but does not eliminate it.

D. Supplementary layers

Three further analyses add texture. When decoupling does occur it is mostly explicit, the ethical issue named rather than implied. Among engaged posts the target of blame shifts with proximity to the harm: artists spread blame across the technology, the companies, the developers and even fellow users, while the smaller body of engaged marketing posts aims more narrowly at companies. Authenticity language, the vocabulary of soul, craft and the human touch, is about twice as common in the creative communities (around 7%) as in marketing (3.5%), and is rarest among decoupling posts (2.8%), which fits the theory: a person who has set the objection aside is unlikely to dwell on what was lost.

E. Limitations in full

Several limitations bound these conclusions and should be read alongside them.

The corpus is drawn from three English-language professional communities on a single platform whose norms shape not only how moral reasoning is expressed but whether it is expressed at all. r/ArtistLounge socially sanctions pro-AI positions, while r/marketing treats moralising as off-topic, so the engagement gap may partly reflect platform governance and social desirability rather than private cognition alone. The patterns reported here should not be generalised to consumers at large or to more reactive platforms.

The three communities are treated as points on a gradient of exposure to generative AI, but that gradient is assumed from occupational composition rather than measured independently. They differ in more than involvement: in demographics, moderation regimes, professional incentives and prevailing discussion norms, any of which could contribute to the observed differences.

The communities also differ in size, and although the analysis uses proportions throughout, the smaller copywriting sample contributes fewer morally engaged posts. The within-engaged comparison rests on 511 posts, with marketing and copywriting contributing few of them, so it should be read as a refinement rather than a settled effect.

Agreement in the substantial band supports the reliability of the scheme as applied, but a rare category is more exposed to residual disagreement than a common one: at 1.9% of the corpus, a small proportion of misapplied labels would represent a meaningful share of the decoupling class. We therefore treat the decoupling rate as an approximate rather than an exact quantity. The community comparison rests on the relative ordering rather than the absolute rates, and that ordering is robust to labelling error distributed evenly across communities. Relatedly, inter-annotator agreement is established on the double-coded pilot rather than on the full corpus, so the reported kappa describes the reliability of the scheme rather than of every label in the dataset. Only the two primary dimensions reached adequate agreement; the explicitness, blame-target and authenticity dimensions plateaued below the threshold and are reported descriptively.

Posts are treated as independent observations, but multiple posts may originate from the same author, which could inflate apparent community differences and permit an author’s characteristic phrasing to appear in both training and test splits. We do not have reliable author-level identifiers after anonymisation, so we cannot quantify the overlap; author-grouped and temporal splits, and user-clustered or hierarchical tests, would place the classification and inference results on firmer ground.

The comparison set is narrow. A linear model over lexical features and a fine-tuned encoder bracket the conventional range, and the zero-shot comparison reported above covers only the 400-post pilot under two prompt designs, so it establishes that this particular regime fails rather than that no general-purpose model could succeed.

Finally, the labels are cross-sectional snapshots of individual posts and do not track the same users over time, so the study describes the distribution of positions in a discourse rather than the evolution of any individual’s reasoning.

F. Future work in full

The most informative extension is a cross-domain transfer test. Because both model families were trained on the pooled corpus, and because the distinctive-vocabulary analysis places *copyright* among the strongest r/ArtistLounge discriminators, the classifiers may be exploiting community register rather than moral-reasoning structure. Training on one community and testing on the others, with comparable macro-F1 as the criterion, would separate the two directly.

Beyond that, comparative breadth matters. Adding communities at intermediate and higher involvement, and platforms with different norms such as X or dedicated artist forums, would test whether the engagement gap is a stable property of involvement or partly an artefact of Reddit culture. A second direction extends the language model comparison beyond the pilot: evaluating a zero-shot baseline across the full label-stripped test set, and against instruction-tuned and few-shot variants, would establish whether the failure we observe is a property of the task or of the prompt regimes we tried. Third, the supplementary dimensions deserve conceptual rather than procedural attention: their plateau suggests the categories for blame attribution and authenticity need redefinition, perhaps grounded in a closer reading of how these communities actually phrase responsibility, before they can support quantitative claims. Fourth, a sensitivity analysis across the decoupling threshold, in particular the treatment of MIXED-stance posts with deontological reasoning, would establish how stable the 1.9 per cent incidence is. Finally, the agenda-setting implication invites an intervention study: if low-involvement users bypass the ethical question rather than resolving it against themselves, then research on how that question is made salient, and whether salience changes behaviour, would speak more directly to policy and platform design than further work on persuasion alone.

References

- [1] J. Lovato, J. W. Zimmerman, I. Smith, P. Dodds, and J. L. Karson, “Foregrounding artist opinions: A survey study on transparency, ownership, and fairness in AI generative art,” in *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (AIES)*, vol. 7, pp. 905–916, 2024.
- [2] T. Eloundou, S. Manning, P. Mishkin, and D. Rock, “GPTs are GPTs: Labor market impact potential of LLMs,” *Science*, vol. 384, no. 6702, pp. 1306–1308, 2024.

-
- [3] A. Bhattacharjee, J. Z. Berman, and A. Reed, “Tip of the hat, wag of the finger: How moral decoupling enables consumers to admire and admonish,” *Journal of Consumer Research*, vol. 39, no. 6, pp. 1167–1184, 2013.
 - [4] K. Haberstroh, U. R. Orth, S. Hoffmann, and B. Brunk, “Consumer response to unethical corporate behavior: A re-examination and extension of the moral decoupling model,” *Journal of Business Ethics*, vol. 140, no. 1, pp. 161–173, 2017.
 - [5] L. Liu, C. Prentice, and J. Z. Marrington, “Ethics on trial: How generative artificial intelligence misuse fuels brand cancellation,” *Journal of Enterprise Information Management*, 2026. Advance online publication.
 - [6] A. Bandura, “Moral disengagement in the perpetration of inhumanities,” *Personality and Social Psychology Review*, vol. 3, no. 3, pp. 193–209, 1999.
 - [7] S. M. Mohammad, S. Kiritchenko, P. Sobhani, X. Zhu, and C. Cherry, “SemEval-2016 task 6: Detecting stance in tweets,” in *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, pp. 31–41, 2016.
 - [8] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, “RoBERTa: A robustly optimized BERT pretraining approach,” *arXiv preprint arXiv:1907.11692*, 2019.
 - [9] C. Du Plessis, “Ethical requirements for generative AI in brand content creation: A qualitative comparative analysis,” *Frontiers in Communication*, vol. 10, p. 1523077, 2025.
 - [10] J. W. Pennebaker, R. L. Boyd, K. Jordan, and K. Blackburn, “The development and psychometric properties of LIWC2015,” tech. rep., University of Texas at Austin, Austin, TX, 2015.